

ANKIT GUPTA

gupta.v.ankit@gmail.com • <http://www.ankitg.me>

EDUCATION

Harvard University, B.A./M.S. in Computer Science

Cambridge, MA

GPA: 3.94/4.0. Magna Cum Laude with Highest Honors. Inducted into Phi Beta Kappa. May 2017

Coursework includes: Deep Learning (CS 287), Machine Learning (CS 181, CS 281), Distributed Comp. (CS 262), Prob. Algorithms (CS 223), Parallel Comp. (CS 205), Data Struct./Alg. (CS 124), Probability (Stat 110), Real Analysis (Math 25)

- Conducted **deep learning thesis research**. Studying transcription factor binding. Published/presented work at ICML.
- Spring 2017 & Spring 2016 Teaching Fellow for CS 181 (Machine Learning) - Taught section, wrote and graded problem sets, and held office hours. Taught four course-wide review sessions per semester.
- Fall 2016 Teaching Fellow for CS 182 (Artificial Intelligence) - Taught section, graded problem sets, and held office hours

TECHNICAL SKILLS

Programming Languages

Python, C++, C, Java, Go, Typescript, SQL, HTML, CSS, MATLAB

Tools

AWS Cloud Infra, PyTorch, Tensorflow, Ray, Docker, Kubernetes, Git

EXPERIENCE

Reverie Labs

Cambridge, MA

Co-Founder and CTO

Nov 2017 — Present

- Co-Founder of Reverie Labs, pharmaceutical company using computation to drive kinase drug discovery
- Built significant infrastructure for distributed training and inference of machine learning models for drug discovery
- Scaled team to 30 employees, and actively growing engineering across multiple roles.
- Directly oversaw org of 14 engineers spanning infrastructure, machine learning, full-stack, and front-end. Managed and grew several individual contributors from new-grad to staff-level engineers, and developed new managers for subteams as the company grew.
- Led the development of a machine learning research/development org delivering dozens of production models for molecular potency and ADMET properties to therapeutic programs in oncology
- Oversaw the development of a modern multi-region, multi-account secure AWS Cloud environment, enabling massive-scale training of ML models with 1000s of GPUs, and large-scale CPU workflows spanning 100,000s of cores
- Led the development of substantial CI/CD automation to enable a smooth developer experience even as a nimble startup, giving in-house developers a modern Dockerized workflow across all develop types
- Contributed extensively to the design and implementation of an AWS-hosted kubernetes cluster running workloads spanning machine learning training, inference, computational chemistry simulations, and software automations

Vicarious AI

San Francisco Bay Area, CA

Research Engineer (Deep Learning)

July 2017 — Nov 2017

- Worked as research engineer at 50-person ML/robotics company with \$130M+ in funding
- Designed and implemented large-scale deep learning models for visual perception

Harvard University School of Engineering and Applied Sciences

Cambridge, MA

Researcher and Machine Learning Teaching Fellow

Jan 2016 — June 2017

- Short Paper “Dilated Convolutions for Modeling Long-Distance Genomic Dependencies” was accepted to the ICML 2017 Workshop on Computational Biology. Invited to give one of four contributed talks.
- Won Speaker Award and Best Poster Award at ICML 2017 workshop. Talk viewable at <https://youtu.be/HmCecphEvQg>.
- Taught sections, held office hours, and wrote and graded problem sets for Harvard’s undergraduate machine learning courses
- Held 8 course-wide exam review sessions across two semesters with 100+ students in attendance at each

Palantir

New York, NY

Software Engineering Intern

May 2016 — Aug 2016

- Worked as full-stack developer on Palantir’s core Spark-based data analysis and visualization product
- Designed back-end for new data transformation prototype product that is used in production across the company
- Engaged closely with team of engineers to actively develop a product used across dozens of high-impact deployments

Google

Mountain View, CA

Software Engineering Intern

May 2015 — August 2015

- Developed back-end software for Google’s content ad targeting teams, using internal parallel data processing tools
- Used text clustering machine learning models to diagnose the source of misclassifications of advertisements